

# 1.Importing required Libraries

```
In [1]: 1 pip install livelossplot
```

...

```
In [35]: 1 import pandas as pd
2 import numpy as np
3 import matplotlib.pyplot as plt
4 import seaborn as sns
```

```
In [22]: 1 from sklearn.metrics import classification_report, confusion_matrix
```

```
In [24]: 1 from sklearn import metrics
```

```
In [3]: 1 from livelossplot.inputs.keras import PlotLossesCallback
```

```
In [4]: 1 import tensorflow as tf
```

```
In [5]: 1 from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
In [6]: 1 from tensorflow.keras.models import Sequential
2
3 from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense, Input, Dropout
4 from tensorflow.keras.layers import BatchNormalization, Activation, MaxPooling2D
5
6
7 from tensorflow.keras.optimizers import Adam, SGD
8 from tensorflow.keras.callbacks import ModelCheckpoint
```

## 2.Load and Pre-process the dataset

## ► Loading the dataset using the pandas read\_csv function

```
In [26]: 1 train_images = pd.read_csv("train_images.csv", header=None)
          2 train_labels = pd.read_csv("train_labels.csv", header=None)
          3 val_images = pd.read_csv("val_images.csv", header=None)
          4 val_labels = pd.read_csv("val_labels.csv", header=None)
```

```
In [8]: 1 train_images.head()
```

```
Out[8]:
```

	0	1	2	3	4	5	6	7	8	9	...	8182	8183	8184	8185
0	0.631373	0.623529	0.713726	0.705882	0.658824	0.666667	0.654902	0.635294	0.647059	0.705882	...	0.682353	0.611765	0.650980	0.658824
1	0.725490	0.752941	0.749020	0.701961	0.690196	0.721569	0.709804	0.745098	0.654902	0.721569	...	0.721569	0.698039	0.721569	0.686275
2	0.717647	0.701961	0.713726	0.733333	0.705882	0.717647	0.725490	0.682353	0.717647	0.674510	...	0.709804	0.694118	0.705882	0.682353
3	0.705882	0.674510	0.654902	0.678431	0.666667	0.662745	0.678431	0.662745	0.686275	0.686275	...	0.639216	0.662745	0.631373	0.643137
4	0.647059	0.729412	0.701961	0.674510	0.611765	0.698039	0.713726	0.662745	0.701961	0.674510	...	0.639216	0.670588	0.705882	0.674510

5 rows × 8192 columns



```
In [9]: 1 train_labels.head()
```

```
Out[9]:
```

	0	1	2	3
0	1.0	0.0	0.0	0.0
1	1.0	0.0	0.0	0.0
2	1.0	0.0	0.0	0.0
3	1.0	0.0	0.0	0.0
4	1.0	0.0	0.0	0.0

### ► Checking the shape of the training and validation data

```
In [10]: 1 print(train_images.shape, train_labels.shape)
```

```
(3200, 8192) (3200, 4)
```

```
In [11]: 1 print(val_images.shape, val_labels.shape)
```


```
(800, 8192) (800, 4)
```

### ► Reshaping the training and validation images

```
In [12]: 1 x_train = train_images.values.reshape(3200, 64, 128, 1)
2 x_val = val_images.values.reshape(800, 64, 128, 1)
3
4 y_train = train_labels.values
5 y_val = val_labels.values
```

## 3. Visualize the dataset

```
In [13]: 1 plt.figure(0, figsize=(15, 15))
2         for i in range(1,4):
3             plt.subplot(1, 3, i)
4             img = np.squeeze(x_train[np.random.randint(0, x_train.shape[0])])
5             plt.xticks([])
6             plt.yticks([])
7             plt.imshow(img, cmap="gray")
```



## 4.Create Training and Validation Data Generators using Keras ImageDataGenerator function

```
In [14]: 1 datagen_train = ImageDataGenerator(horizontal_flip=True,
2                                         rotation_range=0.2,
3                                         width_shift_range=0.2,
4                                         height_shift_range=0.2)
5         datagen_train.fit(x_train)
6
7         datagen_val = ImageDataGenerator(horizontal_flip=True,
8                                             rotation_range=0.2,
9                                             width_shift_range=0.2,
10                                            height_shift_range=0.2)
11        datagen_val.fit(x_val)
```

## 5.Design a Convolutional Neural Network (CNN) Model

In [15]:

```
1  # INITIALIZING THE CNN
2  model = Sequential()
3
4  # FIRST CONVOLUTION
5  model.add(Conv2D(32, (5,5), padding='same', input_shape=(64, 128, 1)))
6  model.add(BatchNormalization())
7  model.add(Activation('relu'))
8  model.add(MaxPooling2D(pool_size=(2,2)))
9  model.add(Dropout(0.2))
10
11 # 2ND CONVOLUTION
12 model.add(Conv2D(64, (5,5), padding='same'))
13 model.add(BatchNormalization())
14 model.add(Activation('relu'))
15 model.add(MaxPooling2D(pool_size=(2,2)))
16 model.add(Dropout(0.2))
17
18 # FLATTEN
19 model.add(Flatten())
20
21 # FULLY CONNECTED LAYER
22 model.add(Dense(1024))
23 model.add(Activation('relu'))
24 model.add(Dropout(0.5))
25
26 # OUTPUT LAYER
27 model.add(Dense(4, activation='softmax'))
```

**6.Compile the Model using Adam optimizer, categorical\_crossentropy loss function, and accuracy matrix**

```
In [16]: 1 initial_learning_rate = 0.005
          2
          3 lr_schedule = tf.keras.optimizers.schedules.ExponentialDecay(
          4     initial_learning_rate = initial_learning_rate,
          5     decay_steps = 5,
          6     decay_rate = 0.96,
          7     staircase=True
          8 )
          9
         10 optimizer = Adam(learning_rate = lr_schedule)
```

```
In [17]: 1 model.compile(optimizer=optimizer,
          2     loss='categorical_crossentropy',
          3     metrics=['accuracy'])
```

## 7. Print the Model summary

In [18]:

```
1 model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 64, 128, 32)	832
batch_normalization (Batch Normalization)	(None, 64, 128, 32)	128
activation (Activation)	(None, 64, 128, 32)	0
max_pooling2d (MaxPooling2D)	(None, 32, 64, 32)	0
dropout (Dropout)	(None, 32, 64, 32)	0
conv2d_1 (Conv2D)	(None, 32, 64, 64)	51264
batch_normalization_1 (Batch Normalization)	(None, 32, 64, 64)	256
activation_1 (Activation)	(None, 32, 64, 64)	0
max_pooling2d_1 (MaxPooling2D)	(None, 16, 32, 64)	0
dropout_1 (Dropout)	(None, 16, 32, 64)	0
flatten (Flatten)	(None, 32768)	0
dense (Dense)	(None, 1024)	33555456
activation_2 (Activation)	(None, 1024)	0
dropout_2 (Dropout)	(None, 1024)	0
dense_1 (Dense)	(None, 4)	4100

=====

Total params: 33612036 (128.22 MB)  
Trainable params: 33611844 (128.22 MB)





Non-trainable params: 192 (768.00 Byte)

---

## **8.Training the Model with batch\_size = 32 & epochs = 12**

In [27]:

```
1 checkpoint = ModelCheckpoint('model_weights.hhs', monitor='val_loss',
2                             save_weights_only=True, mode='min', verbose = 0)
3
4 callbacks = [PlotLossesCallback(), checkpoint]
5
6 batch_size = 32
7
8 history = model.fit(datagen_train.flow(x_train, y_train, batch_size=batch_size,
9                                       shuffle=True),
10                    steps_per_epoch = len(x_train) // batch_size,
11                    validation_data = datagen_val.flow(x_val, y_val, batch_size=batch_size,
12                                                       shuffle=True),
13                    epochs = 12,
14                    callbacks=callbacks)
```



```

accuracy
      training      (min:    0.691, max:    0.712, cur:    0.698)
      validation    (min:    0.706, max:    0.728, cur:    0.716)
Loss
      training      (min:    0.538, max:    0.666, cur:    0.592)
      validation    (min:    0.475, max:    0.490, cur:    0.480)
100/100 [=====] - 102s 1s/step - loss: 0.5919 - accuracy: 0.6981 - val_loss: 0.4805 - val_a
ccuracy: 0.7163

```

## 9. Evaluate the Model

► Use the `model.evaluate` function to evaluate the accuracy

In [28]:

```
1 model.evaluate(x_val, y_val)
```

```
25/25 [=====] - 3s 104ms/step - loss: 0.5685 - accuracy: 0.6712
```

Out[28]: [0.5685462355613708, 0.6712499856948853]

## ► Print a Classification Report and the accuracy score (classification accuracy)

```
In [29]: 1 y_true = np.argmax(y_val, 1)
          2 y_pred = np.argmax(model.predict(x_val), 1)
          3
          4 print(metrics.classification_report(y_true, y_pred))
```

```
25/25 [=====] - 3s 108ms/step
```

	precision	recall	f1-score	support
0	0.63	0.99	0.77	200
1	0.48	0.66	0.55	200
2	0.57	0.04	0.07	200
3	1.00	1.00	1.00	200
accuracy			0.67	800
macro avg	0.67	0.67	0.60	800
weighted avg	0.67	0.67	0.60	800

## ► Display a Confusion Matrix to evaluate the performance of the model

```
In [33]: 1 y_true = np.argmax(y_val, 1)
          2 y_pred = np.argmax(model.predict(x_val), 1)
          3
          4 conf_matrix = metrics.confusion_matrix(y_true, y_pred)
          5 print(metrics.confusion_matrix(y_true, y_pred))
```

```
25/25 [=====] - 3s 105ms/step
```

```
[[198  2  0  0]
 [ 63 131  6  0]
 [ 52 140  8  0]
 [  0  0  0 200]]
```

In [ ]: 1

-----Thank you-----

In [ ]: 1